

RidgeII

February 10, 2019

0.1 Linear Regression: Review

The Boston Housing example.

CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	price
0.00632	18	2.31	0	0.538	6.575	65.2	4.090	1	296		15.3	396.9	4.98
0.02731	0	7.07	0	0.469	6.421	78.9	4.967	2	242		17.8	396.9	9.14
0.02729	0	7.07	0	0.469	7.185	61.1	4.967	2	242		17.8	392.8	4.03
0.03237	0	2.18	0	0.458	6.998	45.8	6.062	3	222		18.7	394.6	2.94
0.06905	0	2.18	0	0.458	7.147	54.2	6.062	3	222		18.7	396.9	5.33

```
In [1]: %matplotlib inline
        from sklearn.datasets import load_boston
        import matplotlib.pyplot as plt
        import numpy as np
        import pandas as pd
        import seaborn as sns
        from pandas.plotting import scatter_matrix
        from sklearn.linear_model import LinearRegression
        import warnings
        from sklearn.preprocessing import StandardScaler

In [2]: warnings.filterwarnings('ignore')

In [3]: boston = load_boston()

In [4]: X = boston.data
        y = boston.target

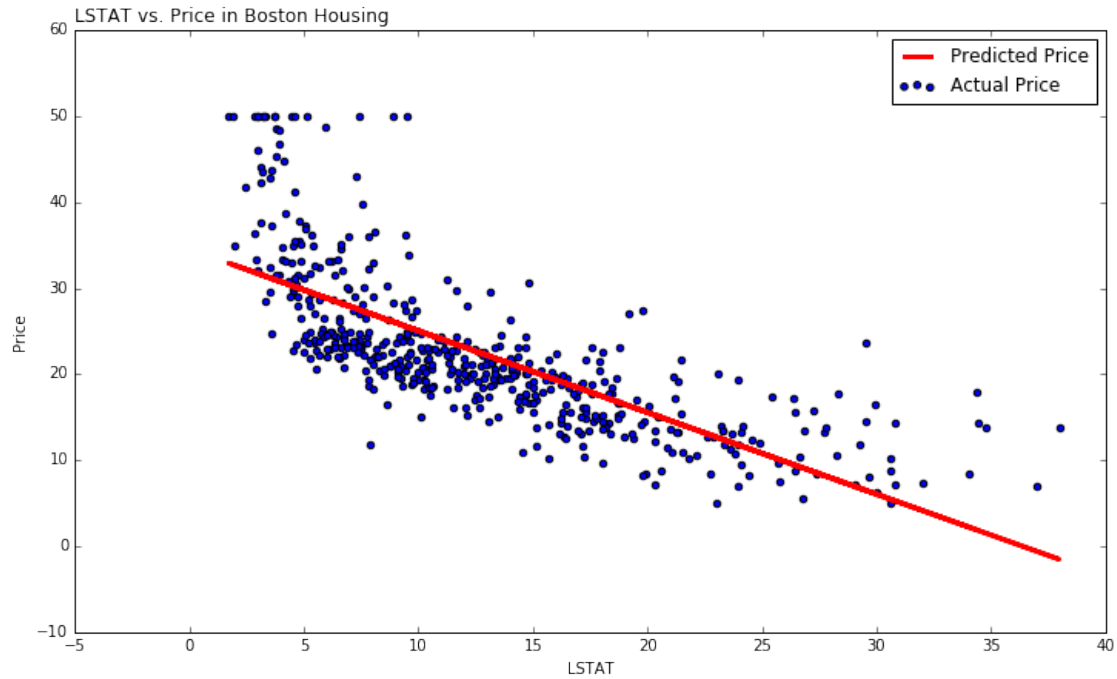
In [5]: bdf = pd.DataFrame(X, columns=boston.feature_names)
        bdf['price'] = y

In [6]: #fitting a scikitlearn model
        lr = LinearRegression()
        lr.fit(bdf[['LSTAT']], y)
        prds= lr.predict(bdf[['LSTAT']])
```

```

In [7]: #making a plot
plt.figure(figsize = (12, 7))
plt.scatter(bdf.LSTAT, bdf.price, label = 'Actual Price')
plt.plot(bdf.LSTAT, prds, color = 'red', linewidth = 3, label = 'Predicted Price')
plt.title('LSTAT vs. Price in Boston Housing', loc = 'left')
plt.xlabel('LSTAT')
plt.ylabel('Price')
plt.legend(frameon = 'false')
plt.savefig('presentation/img1.png')

```



0.2 Model Expression

$$w_{LS} = (X^T X)^{-1} X^T y$$

```

In [8]: #calculate product inside parenthesis
xtx = X.T@X

```

```

In [9]: #find inverse
inv = np.linalg.inv(xtx)

```

```

In [10]: #compute weights
wLS = inv@X.T@y

```

0.3 Writing a Function

take in X and y

compute xtx

compute inverse

compute weights

return weights

```
In [11]: def wLS(X, y):
        '''
        This function provides an
        ordinary least squares fit of a dataset  $X$  on
        a target variable  $y$ .
        ----
         $X$  = input array of feature variables
         $y$  = target array of feature variables
        returns
        array of weights for basic linear regression
        '''
        #check dimensions
        if X.shape[0] < X.shape[1]:
            X = X.T
        #calculate product inside parenthesis
        xtx = X.T@X
        #find inverse
        inv = np.linalg.inv(xtx)
        #compute weights
        prod = inv@X.T
        wLS = prod@y
        return wLS
```

```
In [12]: wLS(X, y)
```

```
Out[12]: array([-9.16297843e-02,  4.86751203e-02, -3.77930006e-03,  2.85636751e+00,
                -2.88077933e+00,  5.92521432e+00, -7.22447929e-03, -9.67995240e-01,
                 1.70443393e-01, -9.38925373e-03, -3.92425680e-01,  1.49832102e-02,
                -4.16972624e-01])
```

0.4 Classes in Python

```
class ClassName:
    <statement-1>
    .
    .
    .
    <statement-N>
```

```

class MyClass:
    """A simple example class"""
    i = 12345

    def f(self):
        return 'hello world'

In [13]: class MyClass:
        """A simple example class"""
        i = 12345

        def f(self):
            return 'hello world'
        #create an instance
        class1 = MyClass()

In [14]: #use our method
        class1.f()

Out[14]: 'hello world'

```

0.4.1 __init__ and self

From the Python docs:

The instantiation operation (“calling” a class object) creates an empty object. Many classes like to create objects with instances customized to a specific initial state. Therefore a class may define a special method named `init()`, like this:

```

def __init__(self):
    self.data = []

class Complex:
    def __init__(self, realpart, imagpart):
        self.r = realpart
        self.i = imagpart

x = Complex(3.0, -4.5)
x.r, x.i

In [15]: class Complex:
        """
        This is a simple class that will
        return a complex number object.
        """
        def __init__(self, realpart, imagpart):
            self.r = realpart
            self.i = imagpart

        x = Complex(3.0, -4.5)
        x.r, x.i

```

```
Out[15]: (3.0, -4.5)
```

```
In [16]: #a different class instance  
x2 = Complex(2.3, 5.6)
```

```
In [17]: #another class instance  
x2.r, x2.i
```

```
Out[17]: (2.3, 5.6)
```

0.5 Our Regression Class

```
class Regression:  
    '''  
    This class contains basic linear  
    regression capabilities.  
    - OLS fit a linear regression model  
    - make predictions with the model  
    '''  
  
    def __init__(self, coefs_, intercept_):  
        self.coefs_ = None  
        self.intercept = None  
  
    def OLS(self, X, y)  
        ...  
        self.coefs_ = wLS  
        return wLS  
  
    def predict(self, X):  
        return predictions  
  
In [18]: class Regression:  
  
        def __init__(self, fit_intercept = True):  
            self.coefs_ = None  
            self.intercept_ = None  
            self._fit_intercept = fit_intercept  
  
        def OLS(self, X, y):  
            '''  
            This function provides an  
            ordinary least squares fit of a dataset X on  
            a target variable y.  
            ----  
            X = input array of feature variables  
            y = target array of feature variables  
            returns  
            array of weights for basic linear regression
```

```

'''
# Check shapes of input matrices.
if X.shape[0] < X.shape[1]:
    X = X.T
if y.shape[0] < y.shape[1]:
    y = y.T

# Prepend ones to x matrix
if self._fit_intercept:
    ones = np.ones((len(y), 1), dtype=int)
    X = np.concatenate((ones, X), axis=1)
else:
    X = X
# fit the model
xtx = X.T@X
inv = np.linalg.inv(xtx)
w_ls = inv@X.T@y
# add intercepts and coefs
self.intercept_ = w_ls[:1]
self.coefs_ = w_ls[1:]
return w_ls

```

In [19]: *#some test cases*

```
lr = Regression()
```

In [20]: `X = np.array([[0, 2], [3, 7], [5, 9], [3.4, 6]])`

```
y = np.array([[2.1, 3.2, 4, 5.6]])
```

In [21]: `lr.OLS(X, y)`

```
Out[21]: array([[ 5.3
 2.41150442],
 [-1.4079646 ]])
```

In [22]: *#model without intercept*

```
lr2 = Regression(fit_intercept=False)
```

In [23]: `lr2.OLS(X, y)`

```
Out[23]: array([[ -0.1530117],
 [ 0.6436483]])
```

0.5.1 Predict

$$\hat{y} = \beta_0 + X\beta_i$$

```
def predict(self, X):
    return self.intercept_ + X@self.coef_
```

In [24]: `del(Regression)`

```
In [25]: class Regression:
```

```
    def __init__(self, fit_intercept = True):
        self.coefs_ = None
        self.intercept_ = None
        self._fit_intercept = fit_intercept

    def wLS(self, X, y):
        '''
        This function provides an
        ordinary least squares fit of a dataset X on
        a target variable y.
        ----
        X = input array of feature variables
        y = target array of feature variables
        returns
        array of weights for basic linear regression
        '''
        # Check shapes of input matrices.
        if X.shape[0] < X.shape[1]:
            X = X.T
        try:
            if y.shape[0] < y.shape[1]:
                y = y.T
        except:
            pass
        # Prepend ones to x matrix
        if self._fit_intercept:
            ones = np.ones((len(y), 1), dtype=int)
            X = np.concatenate((ones, X), axis=1)
        else:
            X = X
        # fit the model
        xtx = X.T@X
        inv = np.linalg.inv(xtx)
        w_ls = inv@X.T@y
        # add intercepts and coefs
        self.intercept_ = w_ls[0]
        self.coefs_ = w_ls[1:]
        return w_ls

    def predict(self, X):
        return self.intercept_ + X@self.coefs_
```

```
In [26]: del(lr)
```

```
In [27]: lr = Regression()
         lr.wLS(X, y)
```

```
Out[27]: array([[ 5.3          ],
                [ 2.41150442],
                [-1.4079646 ]])
```

```
In [28]: lr.coefs_
```

```
Out[28]: array([[ 2.41150442],
                [-1.4079646 ]])
```

```
In [29]: lr.intercept_
```

```
Out[29]: array([5.3])
```

```
In [30]: lr.predict(X)
```

```
Out[30]: array([[2.4840708 ],
                [2.67876106],
                [4.68584071],
                [5.05132743]])
```

```
In [31]: #testing on bigger data
        X = boston.data
        y = boston.target
```

```
In [32]: bos_reg = Regression()
```

```
In [33]: bos_reg.wLS(X, y)
```

```
Out[33]: array([ 3.64911033e+01, -1.07170557e-01,  4.63952195e-02,  2.08602395e-02,
                2.68856140e+00, -1.77957587e+01,  3.80475246e+00,  7.51061703e-04,
                -1.47575880e+00,  3.05655038e-01, -1.23293463e-02, -9.53463555e-01,
                9.39251272e-03, -5.25466633e-01])
```

0.5.2 Error

Sum of squared error: $\sum_{i=1}^n (\hat{y}_i - y_i)^2$

Total sum of squared error: $\sum_{i=1}^n (\bar{y} - y_i)^2$

$$r^2 = 1 - \frac{sse}{tss}$$

Mean Squared Error: $\frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$

```
In [34]: def r2(actual_y, predicted_y):
        sse = np.sum((predicted_y - actual_y)**2)
        tse = np.sum((actual_y - np.mean(actual_y))**2)
        return 1 - sse/tse
```

```
In [35]: def mse(actual_y, predicted_y):
        return np.mean((actual_y - predicted_y)**2)
```



```
In [36]: class Metrics:
```

```
    def __init__(self, X, y, model):
        self.data = X
        self.target = y
        self.model = model

    def r2(self):
        squared_errors = (self.target - self.model.predict(self.data))**2
        sse = np.sum(squared_errors)
        tse = np.sum((self.target - np.mean(self.target))**2)
        return 1 - sse/tse

    def mse(self):
        return np.mean((self.target - self.model.predict(self.data))**2)

    def rmse(self):
        return self.mse()**0.5

    def summary_printed(self):
        print('The r2 score is {:.4}\n\nThe Mean Squared Error is {:.4}\n\nand the RMSE is {:.4}')
```

```
In [37]: lr = Regression()
         lr.wLS(X, y)
         performance = Metrics(X, y, lr)
```

```
In [38]: performance.summary_printed()
```

```
The r2 score is 0.7406
The Mean Squared Error is 21.9
and the RMSE is 4.68
```

$$y = mx + b$$

0.5.3 Polynomial Features

$$y = a + bx_i + cx_i^2$$

```
In [39]: from sklearn.preprocessing import PolynomialFeatures
         x = np.array([2, 3, 4])
         poly = PolynomialFeatures(3, include_bias=False)
         poly.fit_transform(x[:, None])
```

```
Out[39]: array([[ 2.,  4.,  8.],
                [ 3.,  9., 27.],
                [ 4., 16., 64.]])
```

0.5.4 sklearn pipelines

```
In [40]: from sklearn.pipeline import make_pipeline
poly_model = make_pipeline(PolynomialFeatures(7),
                           LinearRegression())
```

```
In [41]: from sklearn.preprocessing import PolynomialFeatures
x = np.array([2, 3, 4])
poly = PolynomialFeatures(3, include_bias=False)
poly.fit_transform(x[:, None])
```


```
Out[41]: array([[ 2.,  4.,  8.],
               [ 3.,  9., 27.],
               [ 4., 16., 64.]])
```

```
In [42]: from sklearn.pipeline import make_pipeline
poly_model = make_pipeline(PolynomialFeatures(7),
                           LinearRegression())
```

```
In [43]: #fitting a high order polynomial to sin with noise
rng = np.random.RandomState(1)
x = 10 * rng.rand(50)
y = np.sin(x) + 0.1 * rng.randn(50)
xfit = np.linspace(0, 10, 1000)

poly_model.fit(x[:, np.newaxis], y)
yfit = poly_model.predict(xfit[:, np.newaxis])

plt.scatter(x, y)
plt.plot(xfit, yfit);
```



0.5.5 Any Basis

```
In [44]: from sklearn.base import BaseEstimator, TransformerMixin

#a class for generating features from gaussian basis
class GaussianFeatures(BaseEstimator, TransformerMixin):
    """Uniformly spaced Gaussian features for one-dimensional input"""

    def __init__(self, N, width_factor=2.0):
        self.N = N
        self.width_factor = width_factor

    @staticmethod
    def _gauss_basis(x, y, width, axis=None):
        arg = (x - y) / width
        return np.exp(-0.5 * np.sum(arg ** 2, axis))

    def fit(self, X, y=None):
        # create N centers spread along the data range
        self.centers_ = np.linspace(X.min(), X.max(), self.N)
        self.width_ = self.width_factor * (self.centers_[1] - self.centers_[0])
        return self

    def transform(self, X):
        return self._gauss_basis(X[:, :, np.newaxis], self.centers_,
                                  self.width_, axis=1)

gauss_model = make_pipeline(GaussianFeatures(20),
                             LinearRegression())
gauss_model.fit(x[:, np.newaxis], y)
yfit = gauss_model.predict(xfit[:, np.newaxis])

plt.scatter(x, y)
plt.plot(xfit, yfit)
plt.xlim(0, 10);
```